

INTRODUCTION

What is a RNN (Recurrent Neural Network) ?

- There are 3 stages of RNN computation, similar to other multi-layer neural networks:
 - the forward pass
 - the backward pass
 - computing gradients during backpropagation

Why use RNNs?

- RNNs are distinguished by **use of previously-computed neuron outputs** across one more axis (not just going "up the layers")
 - RNNs are particularly effective for predictive challenges on data that occur on a segmentable continuum (e.g. time)
 - each segment has intuitive dependencies (that the RNN aims to capture quantitatively) on neighboring segments

We focus on the inference-time forward pass

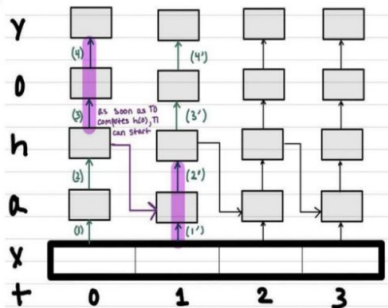
- assume that the training component has completed with parameter values (that are ready for use in inference computations)
- entails the **below four core formulas for layers' neuronal outputs** that we examine for opportunities to parallelize:

$$\begin{aligned} a[t] &= b + W * h[t-1] + U * x[t] \\ h[t] &= \tanh(a[t]) \\ o[t] &= c + V * h[t] \\ y[t] &= \text{softmax}(o[t]) \end{aligned}$$

a, h, and o are (vectors of) intermediate layers' hidden values.

Parallelizing Inference:

- thread for one timestep has to wait for the value of $h[t-1]$ to get $a[t]$.
- project's premise is to mitigate parallelization hindrance of this feature of RNNs



Actions highlighted in purple can occur in parallel via different threads

CUDA Implementation: Challenges and Iterations

CHALLENGE 1: very slow memory access

- Our original implementation used global memory:
 - many issues with this: one was undefined behavior caused by race conditions (see diagram on below-right)
 - Global memory is much slower than shared memory (15x+ according to NVIDIA developer guide docs)

SOLUTION 1: maximally use shared memory local to kernels

- Allocated within a kernel launch
 - Added perk: faster declaration/initialization

CHALLENGE 2: race conditions in writes and reads

- Pausing thread (reading) might collide with thread that is writing -> problems!

SOLUTION 2: Use CUDA primitives for atomicity guarantees

- ATOMIC READS:** use `atomicCAS`
- ATOMIC WRITES:** use `atomicExch`

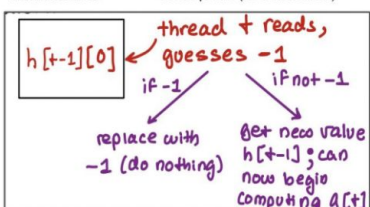
atomicExch

For all timesteps' h vectors,

```
atomicExch(&h[0],
  neuron_output)
atomicExch(&h[1],
  neuron_output)
.
.
atomicExch(&h[HSIZE-1],
  neuron_output)
```

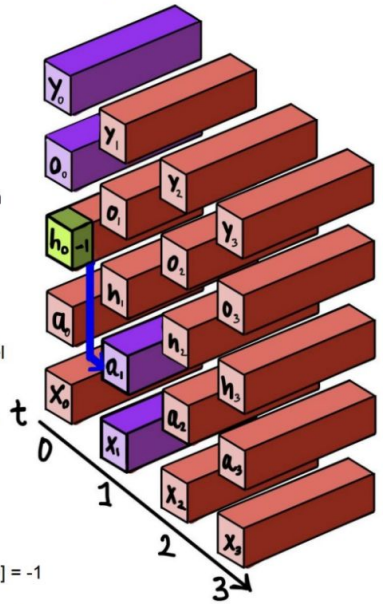
atomicCAS

Note: The CAS provided by CUDA API only works for integers: we had to adapt this (see references)



CUDA Implementation: Logic and Initial Design Decisions

- The CUDA implementation was implemented to analyze the results of a **GPU-based acceleration**.
- We implemented two versions:
 - A sequential version:** this is used as a benchmark to determine the speedup that results from the parallel version
 - A parallelized version**
- The Kernel Function:
 - Sequential version: calls the kernel function with one block with one thread
 - Parallel version: calls the kernel function with one block with *timesteps* number of threads (each thread handles one time step)



OUR DESIGN (see diagram above):

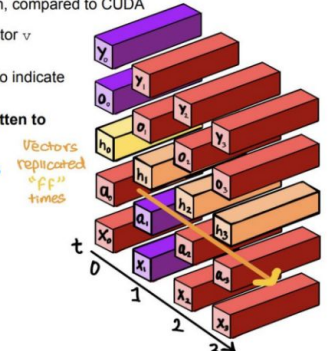
- For each $h[t]$ at each time step: $h[0] = -1$ (labeled in green in the diagram)
 - Thread 1 handles computation for $t = 1$
 - Must wait for Thread 0 to finish computing $h[0]$ before thread 1 can compute $a[1]$
 - Once Thread 0 computes $h[0]$, Thread 1 begins computation for its timestep
 - In purple is the computation that is done in parallel (Thread 0 continues computing $o[0]$ and $y[0]$ while Thread 1 computes $a[1]$)
 - The blue arrow highlights the dependency between threads
- The kernel function
 - Contains spin loop that keeps reading value in $h[t-1]$

OpenMP Implementation: Idea

- The OpenMP implementation is composed in CPU code
 - more tractable with logic involving more branching and variation.
 - So, **outsource the mainloop iteration** (that computes the forward pass) to a function with internally-dependent control flow (e.g. conditionals)
 - but don't have to work on minimization of non-uniformity in that function's execution pattern, compared to CUDA
- Correctness ensured by tracking a vector v
 - initialized to all 0s at timestep 0
 - Previous timestep sets flag to 1 to indicate availability of neuron output
 - v atomically read from and written to

OpenMP Implementation: Challenges

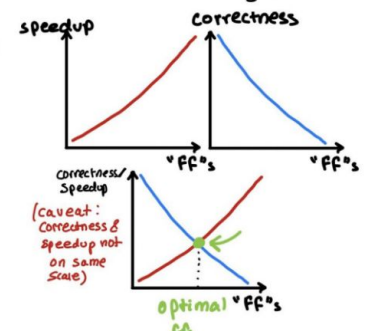
- cost of threads remaining idle observed to be higher with the OpenMP impl.
- Mitigation
 - compromise some accuracy for speedup
 - by way of the strategy we call "forward-fill"



"Forward Fill"

- once $h[t]$ generated by thread at time t , get average ($h[t]$)
- vector of average ($h[t]$) replicated across a tunable parameter of (" ff ") timesteps' hidden neuron outputs
 - tune to find the optimal point tradeoff between degraded correctness and speedup

Tuning result: found that selecting a $ff \approx \text{num_of_threads} + 2$ optimal



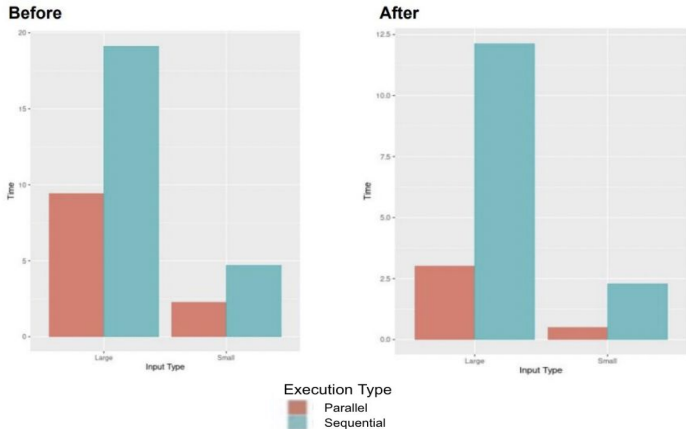
Results for CUDA Implementation

These are the parameters for the two types of problems on which the implementation was tested

Problem Type / Parameters	VSIZE	HSIZE	TIMESTEPS
Small Input Size	8000	50	5000
Large Input Size	8000	125	10000

- TIMESTEPS: number of units of time (typically seconds) that the data spans
- HSIZE: number of features computed in the hidden state of the network (i.e. the dimensionality of the hidden vector h for each time step)
- VSIZE: number of values in the output vector (over which an argmax will yield the prediction)
- For all tests, the number of threads launched per block was 5 (empirically found to be a good balance between (under)-utilization and speedup)

Speedup Plots: Before-and-After switch to per-kernel shared memory access (for h storage)



Discussion of Discoveries, Shortcomings, and Matters for Further Inquiry

CUDA:

- Using shared memory (per-kernel) was very conducive to speedup
- But we still have some aggregation of results that involves global memory
 - May be some optimizations in that approach we overlooked
 - For sufficiently small problem parameters, don't bother with global memory at all?
 - Compute fragmented sub-solutions (of contiguous sections of data) in separate executions and combine later, for bigger input sizes?

Handwritten notes:

requires global memory

6 pay later

INSTEAD

subproblem 0

subproblem 1

subproblem 2

use final subproblem results to reconstruct

use final subproblem results to reconstruct

- OpenMP**
- Forward-filling turned out to be great for speedup
 - However....
 - Its underlying logic is somewhat unsophisticated
 - interpolating data by copying an average, pretty much
 - Did not impinge correctness too much when the ff exceeded $num_of_threads$ by 1-2 (Occam's Razor!), but correctness declined precipitously thereafter
 - Maybe doing some inference on the values to supply in the forward fill can mitigate this!
 - We are doing machine learning for value prediction after all :)
 - But, might require additional compute power at training time, to also develop neuron-like mechanisms for ff 's values

Handwritten notes:

Learning

... (1) network's parameters

AND

... (2) parameters to interpolate network's outputs (?)

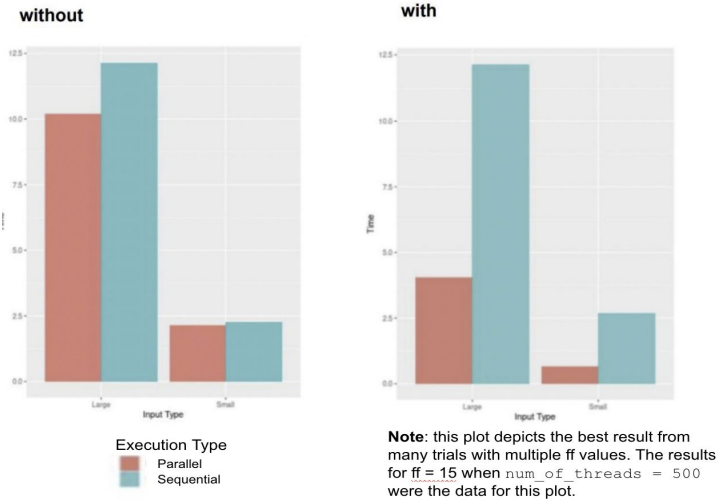
could this help us get more speedup?

We possibly tried this

Results for OpenMP Implementation

- We observed middling speedups with a logical replication of the code we had for CUDA kernels.
 - So, we devised the forward-filling strategy

Speedups with and without Forward-Filling



Because GPUs are much faster than CPUs, we chose the following (smaller) problem parameters to increase comparability with our CUDA impl.

Problem Type / Parameters	VSIZE	HSIZE	TIMESTEPS
Small Input Size	8000	50	500
Large Input Size	8000	125	1000

References

Nabi, J. (2019, July 21). Recurrent neural networks (rnns). Medium. Retrieved December 10, 2021, from <https://towardsdatascience.com/recurrent-neural-networks-rnns-3f06d7653a85>.

Nvidia. (n.d.). Libcudacxx/atomic_thread_fence.MD at main · NVIDIA/libcudacxx. GitHub. Retrieved December 10, 2021, from https://github.com/NVIDIA/libcudacxx/blob/main/docs/extended_api/synchronization_primitives/atomic/atomic_thread_fence.md.

Using shared memory in CUDA C/C++. NVIDIA Developer Blog. (2021, October 29). Retrieved December 10, 2021, from <https://developer.nvidia.com/blog/using-shared-memory-cuda-cc/>.

GPU computing with CUDA Lecture 3 - efficient shared ... - bu. (n.d.). Retrieved December 10, 2021, from <https://www.bu.edu/pasi/files/2011/07/Lecture31.pdf>.

Cuda C++ Programming Guide. NVIDIA Documentation Center. (n.d.). Retrieved December 10, 2021, from <https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html>.

Cuda - tutorial 4 - atomic operations. The Supercomputing Blog. (2011, September 11). Retrieved December 10, 2021, from <http://supercomputingblog.com/cuda/cuda-tutorial-4-atomic-operations/>.

Understanding and using atomic memory operations. Understanding and Using Atomic Memory Operations. (n.d.). Retrieved December 10, 2021, from <https://on-demand.gputechconf.com/gtc/2013/presentations/S3101-Atomic-Memory-Operations.pdf>.

How to access global memory efficiently in CUDA C/C++ kernels. NVIDIA Developer Blog. (2020, August 25). Retrieved December 10, 2021, from <https://developer.nvidia.com/blog/how-access-global-memory-efficiently-cuda-c-kernels>.

Fisseha Berhane, Phd. Building a Recurrent Neural Network - Step by Step - v1. (n.d.). Retrieved December 10, 2021, from https://datascience-enthusiast.com/DL/Building_a_Recurrent_Neural_Network-Step_by_Step_v1.html.