

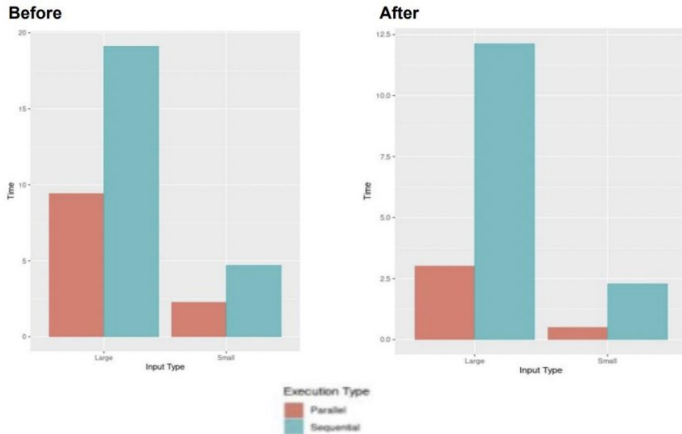
Results for CUDA Implementation

These are the parameters for the two types of problems on which the implementation was tested

Problem Type / Parameters	VSIZE	HSIZE	TIMESTEPS
Small Input Size	8000	50	5000
Large Input Size	8000	125	10000

- TIMESTEPS: number of units of time (typically seconds) that the data spans
- HSIZE: number of features computed in the hidden state of the network (i.e. the dimensionality of the hidden vector h for each time step)
- VSIZE: number of values in the output vector (over which an argmax will yield the prediction)
- For all tests, the number of threads launched per block was 5 (empirically found to be a good balance between (under)-utilization and speedup)

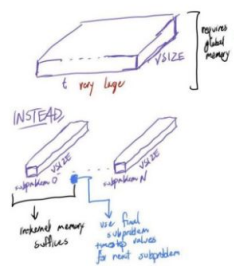
Speedup Plots: Before-and-After switch to per-kernel shared memory access (for h storage)



Discussion of Discoveries, Shortcomings, and Matters for Further Inquiry

CUDA:

- Using shared memory (per-kernel) was very conducive to speedup
- But we still have some aggregation of results that involves global memory
 - May be some optimizations in that approach we overlooked
 - For sufficiently small problem parameters, don't bother with global memory at all?
 - Compute fragmented sub-solutions (of contiguous sections of data) in separate executions and combine later, for bigger input sizes?



OpenMP

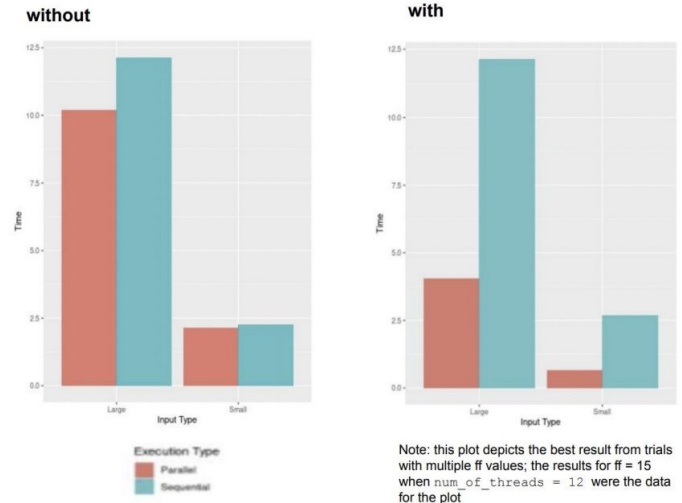
- Forward-fill turned out to be great for speedup
- However....
 - Its underlying logic is somewhat unsophisticated
 - interpolating data by copying an average, pretty much
 - Did not impinge correctness too much when the ff exceeded num_of_threads by 1-2 (Occam's Razor!), but correctness declined precipitously thereafter
 - Maybe doing some inference on the values to supply in the forward fill can mitigate this!
 - We are doing machine learning for value prediction after all :)
 - But, might require additional compute power at training time, to also develop neuron-like mechanisms for ff's values

Learning
 ... (1) network's parameters
 AND
 ... (2) parameters to interpolate network's outputs (?)
 could this help us get more speedup?

Results for OpenMP Implementation

- We observed middling speedups with a logical replication of the code we had for CUDA kernels.
 - So, we devised the forward-filling strategy

Speedups with and without Forward-Filling



Because GPUs are much faster than CPUs, we chose the following (smaller) problem parameters to increase comparability with our CUDA impl.

Problem Type / Parameters	VSIZE	HSIZE	TIMESTEPS
Small Input Size	8000	50	500
Large Input Size	8000	125	1000

References

- Nabi, J. (2019, July 21). Recurrent neural networks (rnns). Medium. Retrieved December 10, 2021, from <https://towardsdatascience.com/recurrent-neural-networks-rnns-3f06d7653a85>.
- Nvidia. (n.d.). Libcudacxx/atomic_thread_fence.MD at main · NVIDIA/libcudacxx. GitHub. Retrieved December 10, 2021, from https://github.com/NVIDIA/libcudacxx/blob/main/docs/extended_api/synchronization_primitives/atomic/atomic_thread_fence.md.
- Using shared memory in CUDA C/C++. NVIDIA Developer Blog. (2021, October 29). Retrieved December 10, 2021, from <https://developer.nvidia.com/blog/using-shared-memory-cuda-cc/>.
- GPU computing with CUDA Lecture 3 - efficient shared ... - bu. (n.d.). Retrieved December 10, 2021, from <https://www.bu.edu/pasi/files/2011/07/Lecture31.pdf>.
- Cuda C++ Programming Guide. NVIDIA Documentation Center. (n.d.). Retrieved December 10, 2021, from <https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html>.
- Cuda - tutorial 4 - atomic operations. The Supercomputing Blog. (2011, September 11). Retrieved December 10, 2021, from <http://supercomputingblog.com/cuda/cuda-tutorial-4-atomic-operations/>.
- Understanding and using atomic memory operations. Understanding and Using Atomic Memory Operations. (n.d.). Retrieved December 10, 2021, from <https://on-demand.gputechconf.com/gtc/2013/presentations/S3101-Atomic-Memory-Operations.pdf>.
- How to access global memory efficiently in CUDA C/C++ kernels. NVIDIA Developer Blog. (2020, August 25). Retrieved December 10, 2021, from <https://developer.nvidia.com/blog/how-access-global-memory-efficiently-cuda-c-kernels>.
- Fisseha Berhane, Phd. Building a Recurrent Neural Network - Step by Step - v1. (n.d.). Retrieved December 10, 2021, from https://datascience-enthusiast.com/DL/Building_a_Recurrent_Neural_Network-Step_by_Step_v1.html.